

Gradient Descent and Linear Regression

Notation

We will be using a lot of index in this lecture so let me spell it out for you. We need two indices: one for telling us which data points and the other for telling which feature I'm talking about. The data points will be denote with upper index in parenthesis. Specifically $\mathbf{x}^{(1)}$ means the vector feature of the first data point, $\vec{x}^{(2)}$ means the feature vector of the second data point and so on. If I want to refer to specific feature I'll be using the lower index. Specifically, $x_1^{(2)}$ means the first feature of the second data point.

In summary, $x_j^{(i)}$ is the j -th feature of the i -th data point.

Linear Regression

Suppose you have the following dataset of the price of a condo and the area. See exercise.

It's clear that you can predict the value of the new data point by just fitting a straight line to the dataset then use the straight line to predict the price of the new condo.

This can be done simply by finding the line that minimize the sum of the distance between the point and the line squared.

$$\text{cost}(h_{\mathbf{w}}) = \frac{1}{2N} \sum_{i=1}^N [h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}]^2$$

where

$$h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

To understand this expression, first, the cost is a function of our hypothesis $h_{\mathbf{w}}$ (line) which is our guess of the true model. More specifically, since the hypothesis is parametrize by the weight the cost function is just a function of weights, \mathbf{w} . Each hypothesis/line has associated cost and we

construct the cost function such that the “best” line is the line that minimize this cost function. Try drawing a few lines and you will see this align with our rough idea of “best” line. The factor of $\frac{1}{2N}$ is just a constant convention; you can ignore it.

The main ingredient is the expression inside the sum over all the data points: each term inside the sum over the data points is the difference between the predicted from our hypothesis($h_{\mathbf{w}}(\mathbf{x}^{(i)})$) value and the true value $y^{(i)}$. Then we square it so that the negative contribution doesn't cancel with positive contribution. The higher the value the further away the line from the point.

The hypothesis, $h_{\mathbf{w}}(\mathbf{x})$, is simply just the linear combination of feature vector, \mathbf{x} . Note also that here we use the convention that $x_0 = 1$. This represents a straight line in 2D and a plane in 3D and hyper plane in higher dimensional data similar to the perceptron. Each \mathbf{w} corresponds to each line.

So all we need to do now is to find \mathbf{w} that minimize the cost function. There is actually a closed form for the solution which is more efficient that what we are going to do here. Those of you taking numerical method before “should” be able to derive the formula from first principle but in this class we are going to use (less efficient) gradient descent to do it.

Gradient Descent

So we have reduced the problem for fitting a line to the problem of just finding parameters, \mathbf{w} , that minimize the cost function. This is a very common task in pattern recognition. We need to pick the best hypothesis, g , from hypothesis set \mathcal{H} . The “best” one has to be special in some way: it has to minimize something.

The gradient descent algorithm is the follow-

74 ing:

75 1. Start at some location.

$$\mathbf{w}' = \mathbf{w} - \eta \nabla f \quad (2)$$

76 2. Look around and walk in the direction
77 which decrease the function the most. This
78 happens to be the gradient, $-\nabla f$

Let me summarize what to do for gradient descent:

111 1. Start with some \mathbf{w}

112 2. Update \mathbf{w} with

$$\mathbf{w}' = \mathbf{w} + \text{step size} \times \left(\frac{-\nabla f}{\|\nabla f\|} \right) \quad (1)$$

$$\mathbf{w}' = \mathbf{w} - \eta \nabla f \quad (3)$$

79 You may be confused with the gradient.
80 The cost function depends only on the
81 weight, the gradient that we are trying to
82 find is of those derivative with respect to
83 each w_i .

Or simultaneously update

$$w'_i = w_i - \eta \frac{\partial}{\partial w_i} f \quad (4)$$

84 3. Keep doing the above until all direction
85 around you only make the function value
86 greater.

It is important to note that the gradient must be computed with old \mathbf{w}

87 See the board for 2D example.

88 There is a question of what we should pick for
89 the step size. We need two things.

116 3. Stop when the the ∇f is low

90 • We want the step size to be large when we
91 are far away from the minimum. This way
92 we can get the minimum quicker.

93 • We want the step size to be small when we
94 are closed to the minimum. This way we
95 won't step over the minimum and we got a
96 good accuracy.

117 This algorithm will be used a lot in this
118 course. You should write your own for once then
119 for the rest just use the package to do this for you.
120 It is simple but the bad news is that this method
121 is susceptible to local minima. This is something
122 to keep in mind when you apply gradient descent
123 to real world application. Fortunately, for the
124 problem of linear regression we are dealing with,
125 the cost function is convex function which guar-
126 antees that there is no local minima.

97 This is sort of recursive question since we
98 want a function that sort of know where the min-
99 imum is to pick the value. However, this is ac-
100 tually an easy thing to do: $\|\nabla f\|$ does what we
101 need. Near the minimum gradient is small(zero
102 small). So we can establish that step size \propto
103 $\|\nabla f\|$. In other words,

127 One can play with the learning rate to get a
128 reasonably fast and accurate convergence. High
129 η is susceptible to going over the minima and may
130 not even converge, but if it does, it will converge
131 very fast. Low η is less susceptible to divergent
132 but it converge slowly. Typical value used is 0.1,
133 0.01 etc.

$$\text{step size} = \eta \|\nabla f\|,$$

104 where η is called the learning rate does ex-
105 actly what we need. The learning rate is typi-
106 cally less than one: 0.01 is a common value. If
107 we plug this into Equation 1, we get a very nice
108 equation:

134 Linear Regression II

135 Now we are equipped with the tools necessary to
136 find the best line to fit it. Let me remind you that
137 this not usually not the way people normally do
138 linear regression but it is a very nice illustration
139 for gradient descent.

140 So, we need to use Equation 4 to update our
 141 weight at each iteration. We need to compute
 142 the gradient. We have two choice of doing this:
 143 numerically or analytically. For this problem, it
 144 is easy enough to do it analytically. So, let us do
 145 it.

$$\begin{aligned}
 \nabla \text{cost} &= \nabla \frac{1}{2N} \sum_{i=1}^N [h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}]^2 \\
 &= \frac{1}{2N} \sum_{i=1}^N \nabla [h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}]^2 \\
 &= \frac{1}{N} \sum_{i=1}^N [h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}] \times \nabla h_{\mathbf{w}}(\mathbf{x}^{(i)}) \\
 &= \frac{1}{N} \sum_{i=1}^N [h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}] \times \mathbf{x}^{(i)} \quad (5)
 \end{aligned}$$

146 The notation may look a bit confusing
 147 putting so let me put back in the index for the
 148 gradient explicitly

$$\frac{\partial}{\partial w_j} \text{cost} = \frac{1}{N} \sum_{i=1}^N [h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}] \times x_j^{(i)} \quad (6)$$

Thus, the update formula for the \mathbf{w} is

$$w'_j = w_j - \eta \frac{\partial}{\partial w_j} \text{cost} \quad (7)$$

150 Let me repeat again that the gradient must
 151 be evaluated at old weight. This means that you
 152 should not overwrite \mathbf{w} while computing the gra-
 153 dient.

154 Take Home Lessons

155 For this lesson I need you to get the big picture.
 156 The fact that you define some sort of cost func-
 157 tion such that the minimization of the cost func-
 158 tion will give you the hypothesis. Plus, the prob-
 159 lem of minimization can be solved with gradient
 160 descent algorithm. You will be doing this over
 161 and over in this course: build a cost function
 162 and minimize it.